

final_project_obesity

April 7, 2024

0.0.1 For this project, I used the obesity dataset from Kaggle:
<https://www.kaggle.com/datasets/aravindpcoder/obesity-or-cvd-risk-classifyregressorcluster>

```
[1]: import pandas as pd

# reading in the obesity dataset
obesity = pd.read_csv("~/obesity.csv")
obesity
```

```
[1]:      Gender      Age      Height      Weight family_history_with_overweight \
0      Female  21.000000  1.620000  64.000000                      yes
1      Female  21.000000  1.520000  56.000000                      yes
2        Male  23.000000  1.800000  77.000000                      yes
3        Male  27.000000  1.800000  87.000000                      no
4        Male  22.000000  1.780000  89.800000                      no
...      ...      ...      ...      ...      ...
2106  Female  20.976842  1.710730  131.408528                      yes
2107  Female  21.982942  1.748584  133.742943                      yes
2108  Female  22.524036  1.752206  133.689352                      yes
2109  Female  24.361936  1.739450  133.346641                      yes
2110  Female  23.664709  1.738836  133.472641                      yes
```

```
      FAVC  FCVC  NCP      CAEC  SMOKE      CH20  SCC      FAF      TUE  \
0      no   2.0  3.0  Sometimes    no  2.000000    no  0.000000  1.000000
1      no   3.0  3.0  Sometimes   yes  3.000000   yes  3.000000  0.000000
2      no   2.0  3.0  Sometimes    no  2.000000    no  2.000000  1.000000
3      no   3.0  3.0  Sometimes    no  2.000000    no  2.000000  0.000000
4      no   2.0  1.0  Sometimes    no  2.000000    no  0.000000  0.000000
...      ...  ...  ...      ...      ...      ...  ...      ...      ...
2106  yes   3.0  3.0  Sometimes    no  1.728139    no  1.676269  0.906247
2107  yes   3.0  3.0  Sometimes    no  2.005130    no  1.341390  0.599270
2108  yes   3.0  3.0  Sometimes    no  2.054193    no  1.414209  0.646288
2109  yes   3.0  3.0  Sometimes    no  2.852339    no  1.139107  0.586035
2110  yes   3.0  3.0  Sometimes    no  2.863513    no  1.026452  0.714137
```

```
      CALC      MTRANS      NObeyesdad
0      no  Public_Transportation  Normal_Weight
```

```

1      Sometimes Public_Transportation Normal_Weight
2      Frequently Public_Transportation Normal_Weight
3      Frequently Walking Overweight_Level_I
4      Sometimes Public_Transportation Overweight_Level_II
...
2106   Sometimes Public_Transportation Obesity_Type_III
2107   Sometimes Public_Transportation Obesity_Type_III
2108   Sometimes Public_Transportation Obesity_Type_III
2109   Sometimes Public_Transportation Obesity_Type_III
2110   Sometimes Public_Transportation Obesity_Type_III

```

[2111 rows x 17 columns]

0.1 Step 1: Data Cleaning

```
[2]: obesity.shape # 2,111 observations and 17 columns
```

```
[2]: (2111, 17)
```

```
[3]: # renaming the columns for convenience
obesity.columns = ['gender', 'age', 'height', 'weight', 'overweight_family',
                  'freq_high_calorie_food', 'freq_vegetable',
                  ↪ 'num_main_meals', 'consume_food_btwn_meals',
                  'smoke', 'daily_h2o_consumption', 'calorie_monitoring',
                  ↪ 'freq_physical_activity',
                  'time_used_technology', 'freq_alcohol_consumption',
                  ↪ 'mode_transportation', 'obesity_level']
```

```
[4]: obesity.head() # checking first few rows of the dataset
```

```
[4]:
```

	gender	age	height	weight	overweight_family	freq_high_calorie_food	\
0	Female	21.0	1.62	64.0	yes	no	
1	Female	21.0	1.52	56.0	yes	no	
2	Male	23.0	1.80	77.0	yes	no	
3	Male	27.0	1.80	87.0	no	no	
4	Male	22.0	1.78	89.8	no	no	

	freq_vegetable	num_main_meals	consume_food_btwn_meals	smoke	\
0	2.0	3.0	Sometimes	no	
1	3.0	3.0	Sometimes	yes	
2	2.0	3.0	Sometimes	no	
3	3.0	3.0	Sometimes	no	
4	2.0	1.0	Sometimes	no	

	daily_h2o_consumption	calorie_monitoring	freq_physical_activity	\
0	2.0	no	0.0	
1	3.0	yes	3.0	

2	2.0	no	2.0
3	2.0	no	2.0
4	2.0	no	0.0

	time_used_technology	freq_alcohol_consumption	mode_transportation \
0	1.0	no	Public_Transportation
1	0.0	Sometimes	Public_Transportation
2	1.0	Frequently	Public_Transportation
3	0.0	Frequently	Walking
4	0.0	Sometimes	Public_Transportation

	obesity_level
0	Normal_Weight
1	Normal_Weight
2	Normal_Weight
3	Overweight_Level_I
4	Overweight_Level_II

```
[5]: # checking dataset for missing values
obesity.isnull().sum()
# There are no missing values
```

```
[5]: gender          0
age                0
height            0
weight            0
overweight_family  0
freq_high_calorie_food  0
freq_vegetable    0
num_main_meals    0
consume_food_btwn_meals  0
smoke             0
daily_h2o_consumption  0
calorie_monitoring  0
freq_physical_activity  0
time_used_technology  0
freq_alcohol_consumption  0
mode_transportation  0
obesity_level      0
dtype: int64
```

```
[6]: # checking all of the levels in the obesity_level variable
obesity['obesity_level'].unique()
```

```
[6]: array(['Normal_Weight', 'Overweight_Level_I', 'Overweight_Level_II',
        'Obesity_Type_I', 'Insufficient_Weight', 'Obesity_Type_II',
        'Obesity_Type_III'], dtype=object)
```

```
[7]: # OBESITY LEVEL VARIABLE
# Here, I define a function to map the levels "Normal Weight" and "Insufficient_
↳Weight" to 0
# and the rest of the levels are mapped to 1
def map_obesity_level(obesity_level):
    if obesity_level in ['Normal_Weight', 'Insufficient_Weight']:
        return 0
    else:
        return 1

# Applying the map_obesity_level function to the 'obesity_level'
# column and create a new column called 'obesity_binary'
obesity['obesity_binary'] = obesity['obesity_level'].apply(map_obesity_level)
```

```
[8]: # OVERWEIGHT FAMILY VARIABLE
# Mapping 'yes' to 1 and 'no' to 0
obesity['overweight_family_binary'] = obesity['overweight_family'].map({'yes': 1,
↳'no': 0})
obesity
```

```
[8]:
```

	gender	age	height	weight	overweight_family	\	
0	Female	21.000000	1.620000	64.000000	yes		
1	Female	21.000000	1.520000	56.000000	yes		
2	Male	23.000000	1.800000	77.000000	yes		
3	Male	27.000000	1.800000	87.000000	no		
4	Male	22.000000	1.780000	89.800000	no		
...		
2106	Female	20.976842	1.710730	131.408528	yes		
2107	Female	21.982942	1.748584	133.742943	yes		
2108	Female	22.524036	1.752206	133.689352	yes		
2109	Female	24.361936	1.739450	133.346641	yes		
2110	Female	23.664709	1.738836	133.472641	yes		
	freq_high_calorie_food		freq_vegetable		num_main_meals		\
0	no		2.0		3.0		
1	no		3.0		3.0		
2	no		2.0		3.0		
3	no		3.0		3.0		
4	no		2.0		1.0		
...		
2106	yes		3.0		3.0		
2107	yes		3.0		3.0		
2108	yes		3.0		3.0		
2109	yes		3.0		3.0		
2110	yes		3.0		3.0		
	consume_food_btwn_meals		smoke		daily_h2o_consumption		calorie_monitoring \

0	Sometimes	no	2.000000	no
1	Sometimes	yes	3.000000	yes
2	Sometimes	no	2.000000	no
3	Sometimes	no	2.000000	no
4	Sometimes	no	2.000000	no
...
2106	Sometimes	no	1.728139	no
2107	Sometimes	no	2.005130	no
2108	Sometimes	no	2.054193	no
2109	Sometimes	no	2.852339	no
2110	Sometimes	no	2.863513	no

	freq_physical_activity	time_used_technology	freq_alcohol_consumption	\
0	0.000000	1.000000		no
1	3.000000	0.000000		Sometimes
2	2.000000	1.000000		Frequently
3	2.000000	0.000000		Frequently
4	0.000000	0.000000		Sometimes
...
2106	1.676269	0.906247		Sometimes
2107	1.341390	0.599270		Sometimes
2108	1.414209	0.646288		Sometimes
2109	1.139107	0.586035		Sometimes
2110	1.026452	0.714137		Sometimes

	mode_transportation	obesity_level	obesity_binary	\
0	Public_Transportation	Normal_Weight	0	
1	Public_Transportation	Normal_Weight	0	
2	Public_Transportation	Normal_Weight	0	
3	Walking	Overweight_Level_I	1	
4	Public_Transportation	Overweight_Level_II	1	
...	
2106	Public_Transportation	Obesity_Type_III	1	
2107	Public_Transportation	Obesity_Type_III	1	
2108	Public_Transportation	Obesity_Type_III	1	
2109	Public_Transportation	Obesity_Type_III	1	
2110	Public_Transportation	Obesity_Type_III	1	

	overweight_family_binary
0	1
1	1
2	1
3	0
4	0
...	...
2106	1
2107	1

```

2108          1
2109          1
2110          1

```

[2111 rows x 19 columns]

```

[9]: # FREQUENCY HIGH CALORIE FOOD
# Mapping 'yes' to 1 and 'no' to 0
obesity['binary_freq_high_calorie_food'] = obesity['freq_high_calorie_food'].
    ↪map({'yes': 1, 'no': 0})
obesity

```

```

[9]:      gender      age      height      weight overweight_family \
0      Female  21.000000  1.620000   64.000000          yes
1      Female  21.000000  1.520000   56.000000          yes
2        Male  23.000000  1.800000   77.000000          yes
3        Male  27.000000  1.800000   87.000000          no
4        Male  22.000000  1.780000   89.800000          no
...      ...      ...      ...      ...      ...
2106  Female  20.976842  1.710730  131.408528          yes
2107  Female  21.982942  1.748584  133.742943          yes
2108  Female  22.524036  1.752206  133.689352          yes
2109  Female  24.361936  1.739450  133.346641          yes
2110  Female  23.664709  1.738836  133.472641          yes

      freq_high_calorie_food  freq_vegetable  num_main_meals \
0                          no              2.0              3.0
1                          no              3.0              3.0
2                          no              2.0              3.0
3                          no              3.0              3.0
4                          no              2.0              1.0
...      ...      ...      ...
2106              yes              3.0              3.0
2107              yes              3.0              3.0
2108              yes              3.0              3.0
2109              yes              3.0              3.0
2110              yes              3.0              3.0

      consume_food_btwn_meals  smoke  daily_h2o_consumption  calorie_monitoring \
0              Sometimes      no              2.000000          no
1              Sometimes     yes              3.000000          yes
2              Sometimes      no              2.000000          no
3              Sometimes      no              2.000000          no
4              Sometimes      no              2.000000          no
...      ...      ...      ...
2106              Sometimes      no              1.728139          no
2107              Sometimes      no              2.005130          no

```

2108	Sometimes	no	2.054193	no
2109	Sometimes	no	2.852339	no
2110	Sometimes	no	2.863513	no

	freq_physical_activity	time_used_technology	freq_alcohol_consumption	\
0	0.000000	1.000000		no
1	3.000000	0.000000		Sometimes
2	2.000000	1.000000		Frequently
3	2.000000	0.000000		Frequently
4	0.000000	0.000000		Sometimes
...
2106	1.676269	0.906247		Sometimes
2107	1.341390	0.599270		Sometimes
2108	1.414209	0.646288		Sometimes
2109	1.139107	0.586035		Sometimes
2110	1.026452	0.714137		Sometimes

	mode_transportation	obesity_level	obesity_binary	\
0	Public_Transportation	Normal_Weight	0	
1	Public_Transportation	Normal_Weight	0	
2	Public_Transportation	Normal_Weight	0	
3	Walking	Overweight_Level_I	1	
4	Public_Transportation	Overweight_Level_II	1	
...	
2106	Public_Transportation	Obesity_Type_III	1	
2107	Public_Transportation	Obesity_Type_III	1	
2108	Public_Transportation	Obesity_Type_III	1	
2109	Public_Transportation	Obesity_Type_III	1	
2110	Public_Transportation	Obesity_Type_III	1	

	overweight_family_binary	binary_freq_high_calorie_food
0	1	0
1	1	0
2	1	0
3	0	0
4	0	0
...
2106	1	1
2107	1	1
2108	1	1
2109	1	1
2110	1	1

[2111 rows x 20 columns]

```
[10]: # CALORIE MONITORING
# Mapping 'yes' to 1 and 'no' to 0
```

```
obesity['binary_calorie_monitoring'] = obesity['calorie_monitoring'].map({'yes':
↪ 1, 'no': 0})
obesity
```

```
[10]:
```

	gender	age	height	weight	overweight_family	\
0	Female	21.000000	1.620000	64.000000		yes
1	Female	21.000000	1.520000	56.000000		yes
2	Male	23.000000	1.800000	77.000000		yes
3	Male	27.000000	1.800000	87.000000		no
4	Male	22.000000	1.780000	89.800000		no
...
2106	Female	20.976842	1.710730	131.408528		yes
2107	Female	21.982942	1.748584	133.742943		yes
2108	Female	22.524036	1.752206	133.689352		yes
2109	Female	24.361936	1.739450	133.346641		yes
2110	Female	23.664709	1.738836	133.472641		yes

	freq_high_calorie_food	freq_vegetable	num_main_meals	\
0	no	2.0	3.0	
1	no	3.0	3.0	
2	no	2.0	3.0	
3	no	3.0	3.0	
4	no	2.0	1.0	
...
2106	yes	3.0	3.0	
2107	yes	3.0	3.0	
2108	yes	3.0	3.0	
2109	yes	3.0	3.0	
2110	yes	3.0	3.0	

	consume_food_btwn_meals	smoke	...	calorie_monitoring	\
0	Sometimes	no	...	no	
1	Sometimes	yes	...	yes	
2	Sometimes	no	...	no	
3	Sometimes	no	...	no	
4	Sometimes	no	...	no	
...
2106	Sometimes	no	...	no	
2107	Sometimes	no	...	no	
2108	Sometimes	no	...	no	
2109	Sometimes	no	...	no	
2110	Sometimes	no	...	no	

	freq_physical_activity	time_used_technology	freq_alcohol_consumption	\
0	0.000000	1.000000	no	
1	3.000000	0.000000	Sometimes	
2	2.000000	1.000000	Frequently	

3	2.000000	0.000000	Frequently
4	0.000000	0.000000	Sometimes
...
2106	1.676269	0.906247	Sometimes
2107	1.341390	0.599270	Sometimes
2108	1.414209	0.646288	Sometimes
2109	1.139107	0.586035	Sometimes
2110	1.026452	0.714137	Sometimes

	mode_transportation	obesity_level	obesity_binary	\
0	Public_Transportation	Normal_Weight	0	
1	Public_Transportation	Normal_Weight	0	
2	Public_Transportation	Normal_Weight	0	
3	Walking	Overweight_Level_I	1	
4	Public_Transportation	Overweight_Level_II	1	
...	
2106	Public_Transportation	Obesity_Type_III	1	
2107	Public_Transportation	Obesity_Type_III	1	
2108	Public_Transportation	Obesity_Type_III	1	
2109	Public_Transportation	Obesity_Type_III	1	
2110	Public_Transportation	Obesity_Type_III	1	

	overweight_family_binary	binary_freq_high_calorie_food	\
0	1	0	
1	1	0	
2	1	0	
3	0	0	
4	0	0	
...	
2106	1	1	
2107	1	1	
2108	1	1	
2109	1	1	
2110	1	1	

	binary_calorie_monitoring
0	0
1	1
2	0
3	0
4	0
...	...
2106	0
2107	0
2108	0
2109	0
2110	0

[2111 rows x 21 columns]

```
[11]: # SMOKE
# Mapping 'yes' to 1 and 'no' to 0
obesity['binary_smoke'] = obesity['smoke'].map({'yes': 1, 'no': 0})
obesity
```

```
[11]:      gender      age      height      weight overweight_family \
0      Female  21.000000  1.620000  64.000000              yes
1      Female  21.000000  1.520000  56.000000              yes
2      Male    23.000000  1.800000  77.000000              yes
3      Male    27.000000  1.800000  87.000000              no
4      Male    22.000000  1.780000  89.800000              no
...
2106  Female  20.976842  1.710730  131.408528              yes
2107  Female  21.982942  1.748584  133.742943              yes
2108  Female  22.524036  1.752206  133.689352              yes
2109  Female  24.361936  1.739450  133.346641              yes
2110  Female  23.664709  1.738836  133.472641              yes

      freq_high_calorie_food  freq_vegetable  num_main_meals \
0                          no              2.0              3.0
1                          no              3.0              3.0
2                          no              2.0              3.0
3                          no              3.0              3.0
4                          no              2.0              1.0
...
2106                      yes              3.0              3.0
2107                      yes              3.0              3.0
2108                      yes              3.0              3.0
2109                      yes              3.0              3.0
2110                      yes              3.0              3.0

      consume_food_btwn_meals  smoke  ...  freq_physical_activity \
0      Sometimes      no  ...              0.000000
1      Sometimes      yes  ...              3.000000
2      Sometimes      no  ...              2.000000
3      Sometimes      no  ...              2.000000
4      Sometimes      no  ...              0.000000
...
2106      Sometimes      no  ...              1.676269
2107      Sometimes      no  ...              1.341390
2108      Sometimes      no  ...              1.414209
2109      Sometimes      no  ...              1.139107
2110      Sometimes      no  ...              1.026452
```

	time_used_technology	freq_alcohol_consumption	mode_transportation \
0	1.000000	no	Public_Transportation
1	0.000000	Sometimes	Public_Transportation
2	1.000000	Frequently	Public_Transportation
3	0.000000	Frequently	Walking
4	0.000000	Sometimes	Public_Transportation
...
2106	0.906247	Sometimes	Public_Transportation
2107	0.599270	Sometimes	Public_Transportation
2108	0.646288	Sometimes	Public_Transportation
2109	0.586035	Sometimes	Public_Transportation
2110	0.714137	Sometimes	Public_Transportation

	obesity_level	obesity_binary	overweight_family_binary \
0	Normal_Weight	0	1
1	Normal_Weight	0	1
2	Normal_Weight	0	1
3	Overweight_Level_I	1	0
4	Overweight_Level_II	1	0
...
2106	Obesity_Type_III	1	1
2107	Obesity_Type_III	1	1
2108	Obesity_Type_III	1	1
2109	Obesity_Type_III	1	1
2110	Obesity_Type_III	1	1

	binary_freq_high_calorie_food	binary_calorie_monitoring	binary_smoke
0	0	0	0
1	0	1	1
2	0	0	0
3	0	0	0
4	0	0	0
...
2106	1	0	0
2107	1	0	0
2108	1	0	0
2109	1	0	0
2110	1	0	0

[2111 rows x 22 columns]

```
[12]: # checking the levels in the consume_food_btwn_meals variable
obesity['consume_food_btwn_meals'].unique()
```

```
[12]: array(['Sometimes', 'Frequently', 'Always', 'no'], dtype=object)
```

```
[13]: # CONSUME FOOD BETWEEN MEALS
# Defining a function to map the categories 'sometimes', 'frequently', and
↳ 'always' to 1
# and everything else to 0
def map_consume_btwn_meals(consumption):
    if consumption in ['Sometimes', 'Frequently', 'Always']:
        return 1
    else:
        return 0

obesity['binary_consume_food_btwn_meals'] = obesity['consume_food_btwn_meals'].
↳ apply(map_consume_btwn_meals)
```

```
[14]: # checking the levels in the freq_alcohol_consumption variable
obesity['freq_alcohol_consumption'].unique()
```

```
[14]: array(['no', 'Sometimes', 'Frequently', 'Always'], dtype=object)
```

```
[15]: # FREQUENT ALCOHOL CONSUMPTION
# Defining a function to map the categories 'sometimes', 'frequently', and
↳ 'always' to 1
# and everything else to 0
def map_consume_alcohol(consume_alcohol):
    if consume_alcohol in ['Sometimes', 'Frequently', 'Always']:
        return 1
    else:
        return 0

obesity['binary_freq_alcohol_consumption'] =
↳ obesity['freq_alcohol_consumption'].apply(map_consume_alcohol)
```

```
[16]: obesity.columns # these are now all of the variables in the dataset
```

```
[16]: Index(['gender', 'age', 'height', 'weight', 'overweight_family',
'freq_high_calorie_food', 'freq_vegetable', 'num_main_meals',
'consume_food_btwn_meals', 'smoke', 'daily_h2o_consumption',
'calorie_monitoring', 'freq_physical_activity', 'time_used_technology',
'freq_alcohol_consumption', 'mode_transportation', 'obesity_level',
'obesity_binary', 'overweight_family_binary',
'binary_freq_high_calorie_food', 'binary_calorie_monitoring',
'binary_smoke', 'binary_consume_food_btwn_meals',
'binary_freq_alcohol_consumption'],
dtype='object')
```

```
[17]: obesity['is_female'] = (obesity['gender'].str.lower() == 'female').astype(int)
obesity
```

```

[17]:      gender      age      height      weight overweight_family \
0      Female  21.000000  1.620000  64.000000          yes
1      Female  21.000000  1.520000  56.000000          yes
2      Male    23.000000  1.800000  77.000000          yes
3      Male    27.000000  1.800000  87.000000          no
4      Male    22.000000  1.780000  89.800000          no
...      ...      ...      ...      ...      ...
2106   Female  20.976842  1.710730  131.408528          yes
2107   Female  21.982942  1.748584  133.742943          yes
2108   Female  22.524036  1.752206  133.689352          yes
2109   Female  24.361936  1.739450  133.346641          yes
2110   Female  23.664709  1.738836  133.472641          yes

      freq_high_calorie_food  freq_vegetable  num_main_meals \
0                          no              2.0              3.0
1                          no              3.0              3.0
2                          no              2.0              3.0
3                          no              3.0              3.0
4                          no              2.0              1.0
...                        ...              ...              ...
2106                        yes            3.0              3.0
2107                        yes            3.0              3.0
2108                        yes            3.0              3.0
2109                        yes            3.0              3.0
2110                        yes            3.0              3.0

      consume_food_btwn_meals  smoke  ...  mode_transportation \
0      Sometimes            no  ...  Public_Transportation
1      Sometimes            yes  ...  Public_Transportation
2      Sometimes            no  ...  Public_Transportation
3      Sometimes            no  ...           Walking
4      Sometimes            no  ...  Public_Transportation
...      ...      ...      ...      ...
2106      Sometimes            no  ...  Public_Transportation
2107      Sometimes            no  ...  Public_Transportation
2108      Sometimes            no  ...  Public_Transportation
2109      Sometimes            no  ...  Public_Transportation
2110      Sometimes            no  ...  Public_Transportation

      obesity_level  obesity_binary  overweight_family_binary \
0      Normal_Weight              0              1
1      Normal_Weight              0              1
2      Normal_Weight              0              1
3      Overweight_Level_I          1              0
4      Overweight_Level_II         1              0
...      ...      ...      ...
2106      Obesity_Type_III          1              1

```

2107	Obesity_Type_III	1	1
2108	Obesity_Type_III	1	1
2109	Obesity_Type_III	1	1
2110	Obesity_Type_III	1	1

	binary_freq_high_calorie_food	binary_calorie_monitoring	binary_smoke \
0	0	0	0
1	0	1	1
2	0	0	0
3	0	0	0
4	0	0	0
...
2106	1	0	0
2107	1	0	0
2108	1	0	0
2109	1	0	0
2110	1	0	0

	binary_consume_food_btwn_meals	binary_freq_alcohol_consumption \
0	1	0
1	1	1
2	1	1
3	1	1
4	1	1
...
2106	1	1
2107	1	1
2108	1	1
2109	1	1
2110	1	1

	is_female
0	1
1	1
2	0
3	0
4	0
...	...
2106	1
2107	1
2108	1
2109	1
2110	1

[2111 rows x 25 columns]

```
[18]: # dropping the mode of transportation variable
obesity.drop(columns='mode_transportation', inplace=True)
```

0.2 Step 2: Data Exploration

```
[19]: import matplotlib.pyplot as plt


weight_obese = obesity.loc[obesity['obesity_binary'] == 1, 'weight']
weight_not_obese = obesity.loc[obesity['obesity_binary'] == 0, 'weight']

plt.hist(weight_obese, color='red', alpha=0.5, label='Obese', bins=20)
plt.hist(weight_not_obese, color='blue', alpha=0.5, label='Not Obese', bins=20)

plt.xlabel('Weight')
plt.ylabel('Frequency')
plt.title('Histogram of Weight Colored by Obesity Level')

plt.legend()

plt.show()
```



Inspecting the Histogram of Weight Colored by Obesity Level, I notice that there is a little bit of overlap between the “obsese” and “not obese” categories, but they are decently separable.

```
[20]: # Histogram of obesity_level variable
obesity_val_counts = obesity['obesity_level'].value_counts()

plt.figure(figsize=(8, 6))

# creating bar plot
plt.bar(obesity_val_counts.index, obesity_val_counts.values, color='blue',
        edgecolor='black', alpha=0.7)

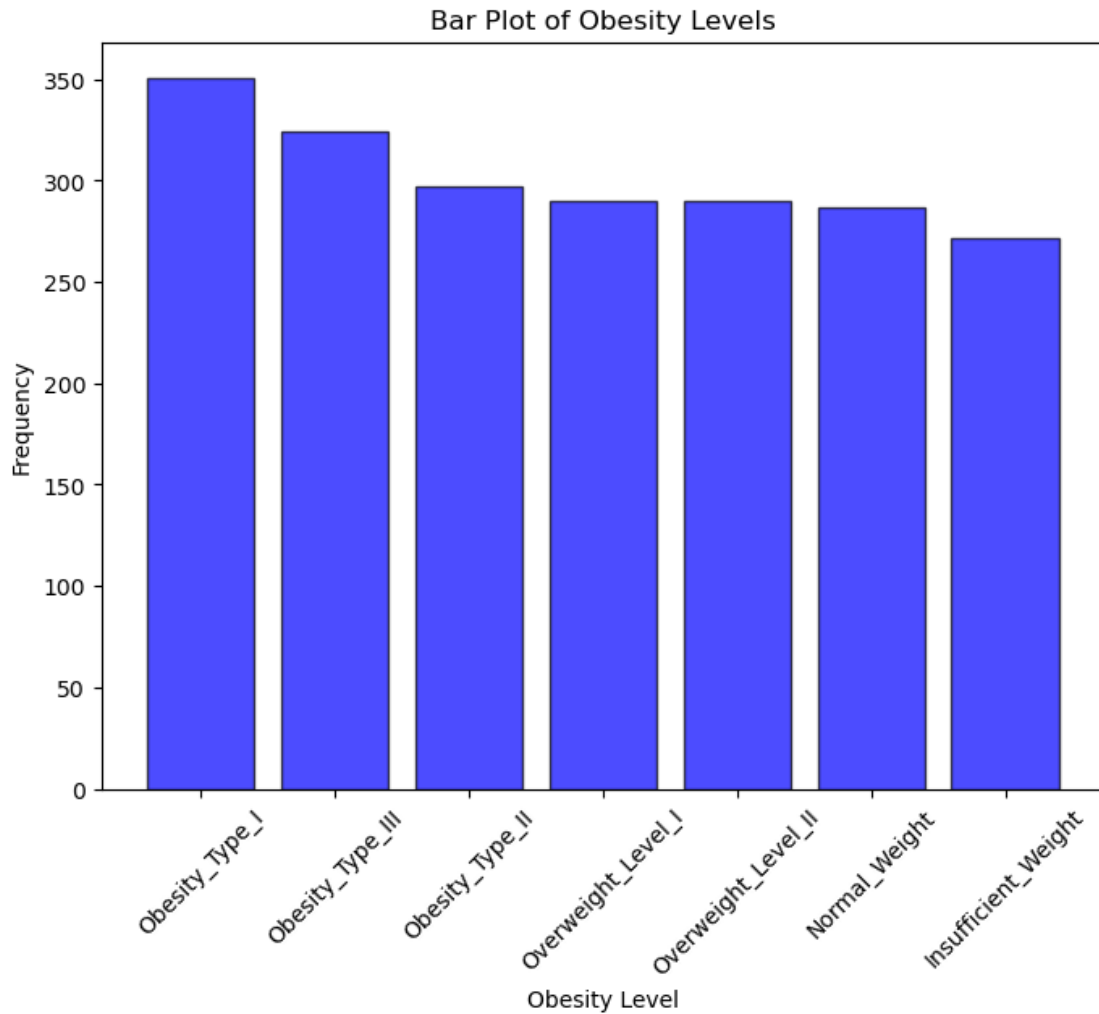
# adding title
plt.title('Bar Plot of Obesity Levels')

# adding x label
plt.xlabel('Obesity Level')

# adding y label
plt.ylabel('Frequency')

# rotating x ticks
plt.xticks(rotation=45)

# Show plot
plt.show()
```

The bar plot of obesity level reveals that the frequencies of each obesity level are not too different from each other.

```
[21]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

X = obesity[['is_female', 'weight', 'age', 'height', 'overweight_family_binary',
↳ 'binary_freq_high_calorie_food',
        'binary_calorie_monitoring', 'binary_smoke', 'freq_vegetable',
        'binary_consume_food_btwn_meals',
↳ 'binary_freq_alcohol_consumption',
        'daily_h2o_consumption', 'num_main_meals',
↳ 'freq_physical_activity', 'time_used_technology']]
y = obesity['obesity_binary']
```


```

data = pd.concat([X, y], axis=1)

correlation_matrix = data.corr()

plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f",
            linewidths=.5)
plt.title('Correlation Matrix for Obesity Dataset')
plt.show()

```



The correlation matrix shows the level of correlation between variables in the dataset. Weight and obesity_binary are strongly, positively correlated. Obesity_binary and overweight_family_binary

are moderately, positively correlated.

0.3 Step 3: Logistic Regression Model

```
[22]: # full model: 15 predictors

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, \
    recall_score, f1_score

X = obesity[['is_female', 'weight', 'age', 'height', 'overweight_family_binary', \
    'binary_freq_high_calorie_food', \
        'binary_calorie_monitoring', 'binary_smoke', 'freq_vegetable', \
        'binary_consume_food_btwn_meals', \
    'binary_freq_alcohol_consumption', \
        'daily_h2o_consumption', 'num_main_meals', \
    'freq_physical_activity', 'time_used_technology']]
y = obesity['obesity_binary']

# 20% of the data goes into training set and rest for testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, \
    random_state = 43)

# initialize logistic regression model
model = LogisticRegression(max_iter = 1000)

# training the model
model.fit(X_train, y_train)

# predict on the testing data
y_pred = model.predict(X_test)

# Find the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# confusion matrix for precision, recall
# Generating confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Calculating precision and recall
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
```

```

print("Confusion Matrix:")
print(cm)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)

```

Accuracy: 0.9550827423167849
 Confusion Matrix:
 [[89 13]
 [6 315]]
 Precision: 0.9603658536585366
 Recall: 0.9813084112149533
 F1 Score: 0.9707241910631741

```

[23]: # Now, using stats models, I can find the pseudo r-squared and p-values
      # to see which variables are most statistically significant.
import statsmodels.api as sm
import numpy as np

X = obesity[['is_female', 'weight', 'age', 'height', 'overweight_family_binary',
             ↪ 'binary_freq_high_calorie_food',
             'binary_calorie_monitoring', 'binary_smoke', 'freq_vegetable',
             'binary_consume_food_btwn_meals',
             ↪ 'binary_freq_alcohol_consumption',
             'daily_h2o_consumption', 'num_main_meals',
             ↪ 'freq_physical_activity', 'time_used_technology']]
y = obesity['obesity_binary']

X_with_const = sm.add_constant(X)

# fitting the model
logit_model = sm.Logit(y, X_with_const)

# getting model results
logit_result = logit_model.fit()

# printing summary
print(logit_result.summary())

```

Optimization terminated successfully.
 Current function value: 0.010834
 Iterations 16

Logit Regression Results			
=====			
Dep. Variable:	obesity_binary	No. Observations:	2111
Model:	Logit	Df Residuals:	2095
Method:	MLE	Df Model:	15
Date:	Thu, 21 Mar 2024	Pseudo R-squ.:	0.9813

```

Time:                22:32:28   Log-Likelihood:        -22.870
converged:            True     LL-Null:                -1220.2
Covariance Type:      nonrobust   LLR p-value:          0.000

```

```

=====
=====

```

		coef	std err	z	P> z
[0.025	0.975]				

const		178.3995	35.808	4.982	0.000
108.218	248.581				
is_female		3.8651	1.717	2.251	0.024
0.500	7.230				
weight		2.7437	0.536	5.115	0.000
1.692	3.795				
age		0.0758	0.064	1.193	0.233
-0.049	0.200				
height		-221.8823	43.272	-5.128	0.000
-306.695	-137.070				
overweight_family_binary		-0.2845	0.988	-0.288	0.774
-2.222	1.653				
binary_freq_high_calorie_food		1.8168	1.232	1.475	0.140
-0.598	4.231				
binary_calorie_monitoring		2.5950	1.334	1.945	0.052
-0.019	5.209				
binary_smoke		-6.3120	10.318	-0.612	0.541
-26.535	13.911				
freq_vegetable		-0.7555	0.851	-0.888	0.374
-2.423	0.912				
binary_consume_food_btwn_meals		0.8899	5.147	0.173	0.863
-9.198	10.978				
binary_freq_alcohol_consumption		0.2359	1.143	0.206	0.837
-2.005	2.477				
daily_h2o_consumption		-0.6487	0.813	-0.797	0.425
-2.243	0.946				
num_main_meals		-0.7463	0.627	-1.190	0.234
-1.976	0.483				
freq_physical_activity		-0.4357	0.467	-0.934	0.350
-1.350	0.479				
time_used_technology		0.5279	0.686	0.769	0.442
-0.817	1.873				

```

=====
=====

```

Possibly complete quasi-separation: A fraction 0.89 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

0.3.1 Suppose we did not have access to weight as a predictor, a.k.a we did not have data on weight. Would we still be able to predict whether or not someone is obese just by looking at the other features?

```
[24]: # Now we can try Logistic Regression without weight as a feature (14 predictors,
      ↪now)

X = obesity[['is_female', 'age', 'height', 'overweight_family_binary',
      ↪'binary_freq_high_calorie_food',
      'binary_calorie_monitoring', 'binary_smoke', 'freq_vegetable',
      'binary_consume_food_btwn_meals',
      ↪'binary_freq_alcohol_consumption',
      'daily_h2o_consumption', 'num_main_meals',
      ↪'freq_physical_activity', 'time_used_technology']]
y = obesity['obesity_binary']

# 20% of the data goes into training set and rest for testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state = 43)

# initialize logistic regression model
model = LogisticRegression(max_iter = 1000)

# training the model
model.fit(X_train, y_train)

# predict on the testing data
y_pred = model.predict(X_test)

# Find the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
# confusion matrix for precision, recall

# Generating confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Calculating precision and recall
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("Confusion Matrix:")
print(cm)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

```
#final logistic model coefficients, p-values. correlation matrix
```

Accuracy: 0.8534278959810875

Confusion Matrix:

```
[[ 57  45]
 [ 17 304]]
```

Precision: 0.8710601719197708

Recall: 0.9470404984423676

F1 Score: 0.9074626865671641

```
[25]: X = obesity[['is_female','age','height','overweight_family_binary',
↳ 'binary_freq_high_calorie_food',
        'binary_calorie_monitoring', 'binary_smoke', 'freq_vegetable',
        'binary_consume_food_btwn_meals'],
↳ 'binary_freq_alcohol_consumption',
        'daily_h2o_consumption','num_main_meals'],
↳ 'freq_physical_activity', 'time_used_technology']]
y = obesity['obesity_binary']

# Add a constant to the features (X) to account for the intercept term
X_with_const = sm.add_constant(X)

# Fit logistic regression model
logit_model = sm.Logit(y, X_with_const)

# Get the fitted model results
logit_result = logit_model.fit()

# Print model summary containing coefficients and p-values
print(logit_result.summary())
```

Optimization terminated successfully.

Current function value: 0.382022

Iterations 7

Logit Regression Results

```
=====
Dep. Variable:          obesity_binary    No. Observations:          2111
Model:                  Logit             Df Residuals:            2096
Method:                 MLE               Df Model:                14
Date:                  Thu, 21 Mar 2024   Pseudo R-squ.:            0.3391
Time:                  22:32:28           Log-Likelihood:           -806.45
converged:              True              LL-Null:                  -1220.2
Covariance Type:        nonrobust         LLR p-value:              1.439e-167
=====
=====
                                coef      std err          z      P>|z|
-----
[0.025      0.975]
```

```

-----
const                -4.4150      1.620      -2.725      0.006
-7.591      -1.239
is_female            -0.1225      0.169      -0.727      0.467
-0.453      0.208
age                  0.1672      0.016      10.372      0.000
0.136      0.199
height              0.1455      1.010      0.144      0.885
-1.834      2.125
overweight_family_binary 2.6148      0.172      15.210      0.000
2.278      2.952
binary_freq_high_calorie_food 0.7317      0.191      3.830      0.000
0.357      1.106
binary_calorie_monitoring 0.0241      0.279      0.086      0.931
-0.523      0.571
binary_smoke        -1.0800      0.440      -2.456      0.014
-1.942      -0.218
freq_vegetable      -0.0942      0.128      -0.737      0.461
-0.345      0.156
binary_consume_food_btwn_meals -0.6130      0.404      -1.519      0.129
-1.404      0.178
binary_freq_alcohol_consumption 0.7606      0.139      5.464      0.000
0.488      1.033
daily_h2o_consumption 0.4745      0.113      4.208      0.000
0.253      0.695
num_main_meals      -0.5013      0.089      -5.658      0.000
-0.675      -0.328
freq_physical_activity -0.3051      0.082      -3.742      0.000
-0.465      -0.145
time_used_technology -0.1316      0.108      -1.221      0.222
-0.343      0.080
=====
=====

```

0.4 Step 4: Decision Tree Classifier Model

```

[26]: # Full model: 15 predictors

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report, \
    mean_squared_error
import pandas as pd

# Split the dataset into training and testing sets

```



```

X_train, X_test, y_train, y_test = train_test_split(
    obesity[['is_female', 'weight', 'age', 'height', 'overweight_family_binary',
            'binary_freq_high_calorie_food', 'freq_vegetable',
            'binary_calorie_monitoring', 'binary_smoke',
            'binary_consume_food_btwn_meals', 'binary_freq_alcohol_consumption',
            'daily_h2o_consumption', 'num_main_meals',
            'freq_physical_activity', 'time_used_technology']],
    obesity['obesity_binary'],
    test_size=0.2, random_state=42)

# Initialize decision tree classifier
tree = DecisionTreeClassifier(random_state=42)

# Train the classifier on the training data
tree.fit(X_train, y_train)

# Predict on the testing data
y_pred = tree.predict(X_test)

# Evaluate the classifier
accuracy = accuracy_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
print("Accuracy:", accuracy)
print("MSE:", mse)

# Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.9716312056737588

MSE: 0.028368794326241134

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.95	0.95	118
1	0.98	0.98	0.98	305
accuracy			0.97	423
macro avg	0.96	0.96	0.96	423
weighted avg	0.97	0.97	0.97	423

```
[27]: # Now trying decision tree without the weight variable (14 predictors now)

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report, \
    mean_squared_error
import pandas as pd

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(obesity[['is_female', 'age', \
    'height', 'overweight_family_binary', \
    'binary_freq_high_calorie_food', 'freq_vegetable', \
    'binary_calorie_monitoring', 'binary_smoke', \
    'binary_consume_food_btwn_meals', 'binary_freq_alcohol_consumption', \
    'daily_h2o_consumption', 'num_main_meals', \
    'freq_physical_activity', 'time_used_technology']], \
    obesity['obesity_binary'], \
    test_size=0.2, random_state=42)

# Initialize decision tree classifier
tree2 = DecisionTreeClassifier(random_state=42)

# Train the classifier on the training data
tree2.fit(X_train, y_train)

# Predict on the testing data
y_pred = tree2.predict(X_test)

# Evaluate the classifier
accuracy = accuracy_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
print("Accuracy:", accuracy)
print("MSE:", mse)

# Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.851063829787234

MSE: 0.14893617021276595

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.74	0.71	0.73	118
1	0.89	0.90	0.90	305
accuracy			0.85	423
macro avg	0.82	0.81	0.81	423
weighted avg	0.85	0.85	0.85	423